


# Time Series Analysis

## Using ARIMA & Prophet


757 PYTHON MEETUP: 13TH, JULY, 2017  
RAJA HARSHA CHINTA (RCHIN001@ODU.EDU)



# What makes a Time Series Special?


- TS is a collection of data points collected at constant time intervals
- Time Dependent
- Consists Trend and Seasonality
- Analyzed to determine the long term trend and forecast the future

	ds	y
0	2007-12-10	9.590761
1	2007-12-11	8.519590
2	2007-12-12	8.183677
3	2007-12-13	8.072467
4	2007-12-14	7.893572





# Stationarity of a Time Series


- TS with a particular behaviour over time  
→ Very high probability to repeat
- The mean of the series should not be a function of time rather should be a constant.
- The variance of the series should not be a function of time.
- The covariance of the  $i$  th term and the  $(i + m)$  th term should not be a function of time.



Non-Stationary series



Non-Stationary series



Stationary series


# Check for Stationarity

## Plotting Rolling Statistics:

- Plot the moving average or moving variance and see if it varies with time.
- More of a Visual technique.

## Dickey-Fuller Test:


- Statistical test where NULL hypothesis is referred as Non-Stationary.
- Comprise of a **Test Statistic**, **Critical Values** for different confidence levels.
- If the '**Test Statistic**' < '**Critical Value**', we can reject the null hypothesis and say that the series is stationary.



```
Results of Dickey-Fuller Test:
Test Statistic      0.815369
p-value             0.991880
#Lags Used          13.000000
Number of Observations Used  130.000000
Critical Value (5%)  -2.884042
Critical Value (1%)  -3.481682
Critical Value (10%) -2.578770
dtype: float64
```

# Stationarize a Time Series

- **Trend & Seasonality** are two main reasons for Non - Stationarity.
- Apply transformations like **Log**, **Square**, **Cube Root**, etc. on TS to Stationarize.
- In real-time scenarios it is inevitable to see noise in data. In these cases we estimate trend by calculating:
  - Monthly, Weekly Averages
  - Rolling Averages (Smoothing)
  - Exponentially Weighted Moving Average (ewma)
  - Fit a regression model




## Results of Dickey-Fuller Test:

Test Statistic	-3.162908
p-value	0.022235
#Lags Used	13.000000
Number of Observations Used	119.000000
Critical Value (5%)	-2.886151
Critical Value (1%)	-3.486535
Critical Value (10%)	-2.579896
dtype:	float64

# Eliminating Trend & Seasonality


Simple Trend reduction techniques don't work in all cases. Particularly, with high seasonality.

1. **Differencing** – Taking the difference with a particular time lag
2. **Decomposition** – Modeling both trend and seasonality and removing them from the model.



```
Results of Dickey-Fuller Test:
Test Statistic      -2.717131
p-value             0.071121
#Lags Used          14.000000
Number of Observations Used  128.000000
Critical Value (5%)  -2.884398
Critical Value (1%)  -3.482501
Critical Value (10%) -2.578960
dtype: float64
```

Differencing



```
Results of Dickey-Fuller Test:
Test Statistic      -6.332387e+00
p-value             2.885059e-08
#Lags Used          9.000000e+00
Number of Observations Used  1.220000e+02
Critical Value (5%)  -2.885538e+00
Critical Value (1%)  -3.485122e+00
Critical Value (10%) -2.579569e+00
dtype: float64
```

Decomposition

# ARIMA - Time Series Forecasting

A series with significant **dependence among values** need to use some statistical models like ARIMA to forecast the data. ARIMA stands for **Auto-Regressive Integrated Moving Averages**.

```
model = ARIMA(ts_log, order=(p, d, q))
```

The predictors depend on the parameters (p,d,q) of the ARIMA model:

- **Number of AR (Auto-Regressive) terms (p):**

AR terms are just lags of dependent variable.

For instance if p is 5, the predictors for  $x(t)$  will be  $x(t-1) \dots x(t-5)$ .

- **Number of MA (Moving Average) terms (q):**

MA terms are lagged forecast errors in prediction equation.

For instance if q is 5, the predictors for  $x(t)$  will be  $e(t-1) \dots e(t-5)$  where  $e(i)$  is the difference between the moving average at  $i$ th instant and actual value.


- **Number of Differences (d):**

These are the number of non-seasonal differences.

# AR & MA Models


- AR Model :  $p \neq 0$  and  $q = 0$

```
model = ARIMA(ts_log, order=(2, 1, 0))
```



- MA Model :  $p = 0$  and  $q \neq 0$

```
model = ARIMA(ts_log, order=(0, 1, 2))
```






# ARIMA Model

Both AR & MA models have RSS values **1.5023**, **1.4721** respectively and are further reduced to **1.0292** after combining both models into ARIMA with order **(2, 1, 2)**.

```
Model = ARIMA(ts_log, order=(2, 1, 2))
```




# Forecasting back in Original Scale

```
model = ARIMA(ts_log, order=(2, 1, 2))
results_ARIMA = model.fit(dispatch=-1)

predictions_ARIMA_diff = pd.Series(results_ARIMA.fittedvalues, copy=True)
predictions_ARIMA_diff_cumsum = predictions_ARIMA_diff.cumsum()

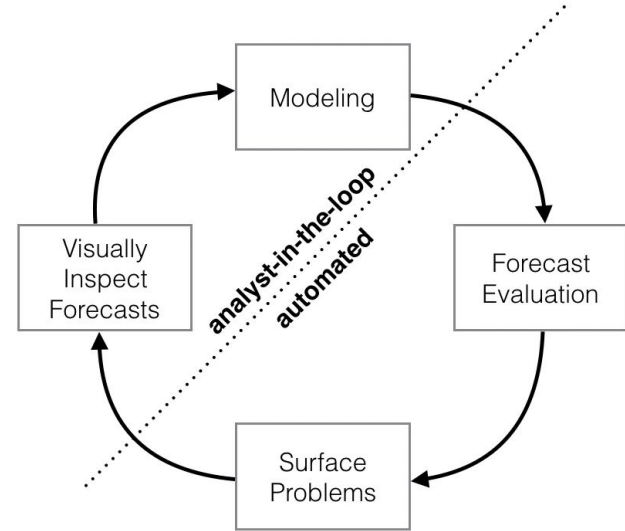
predictions_ARIMA_log = pd.Series(ts_log.ix[0], index=ts_log.index)
predictions_ARIMA_log = predictions_ARIMA_log.add(predictions_ARIMA_diff_cumsum, fill_value=0)

predictions_ARIMA = np.exp(predictions_ARIMA_log)
```



# Facebook's Prophet

- Prophet is a procedure for forecasting time series data.
- It is based on an additive model where non-linear trends are fit with yearly and weekly seasonality, plus holidays.
- It works best with daily periodicity data with at least one year of historical data.
- Prophet is robust to missing data, shifts in the trend, and large outliers.
- Accurate, Fast, Fully Automatic and Tunable forecasts.



# Forecasting using Prophet

- The input to Prophet is always a dataframe with two columns: `ds` and `y`.
- **`ds`** & **`y`** - Represents the date stamp and measurement we wish to forecast.
- We fit the model by instantiating a new `Prophet` object and use `Fit` method.

```
Model = Prophet.fit(df)
```

- Create future dataframe by using default Prophet method - `make_future_dataframe`


```
Future = Model.make_future_dataframe(periods=365)
```

- The predict method will assign each row in future a predicted value which it names `yhat`.

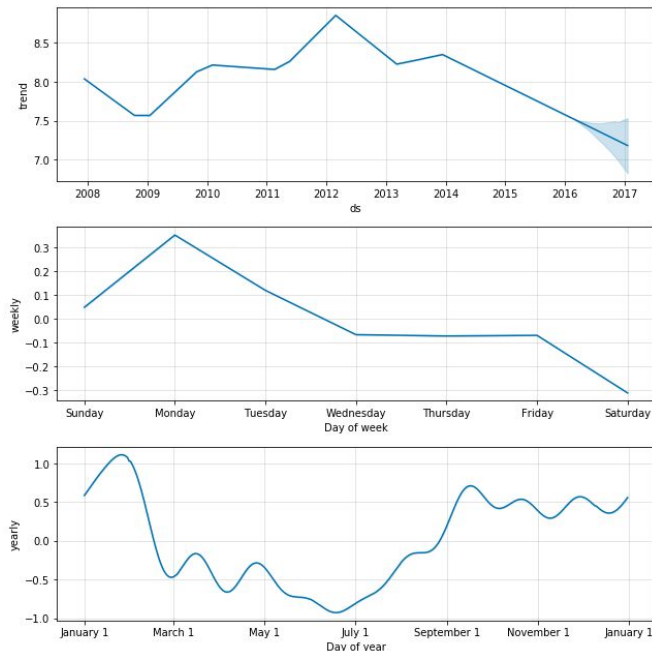
```
Forecast = Model.predict(Future)
```

# Visualizing Model Results

- Pass forecast dataframe and Plot by calling the `Prophet.plot` method.
- Plot forecast components like trend, yearly, weekly seasonality and holidays can be Plot using `Prophet.plot_components` method.



```
m.plot(forecast);
```



```
m.plot_components(forecast);
```



**Thank you**  
Happy Forecasting!